

Importing libraries

```
In [ ]: # Importing all libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import accuracy_score, recall_score, precision_score,

from sklearn.svm import SVC
```

Loading the Dataset

```
In [ ]: path = 'Credit_Card.csv'
df = pd.read_csv(path)
```

Viewing the Dataset

```
In [ ]: df
```

```
Out[ ]:
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4
0	1	20000.0	2	2	1	24	2	2	-1	-1
1	2	120000.0	2	2	2	26	-1	2	0	0
2	3	90000.0	2	2	2	34	0	0	0	0
3	4	50000.0	2	2	1	37	0	0	0	0
4	5	50000.0	1	2	1	57	-1	0	-1	0
...
29995	29996	220000.0	1	3	1	39	0	0	0	0
29996	29997	150000.0	1	3	2	43	-1	-1	-1	-1
29997	29998	30000.0	1	2	2	37	4	3	2	-1
29998	29999	80000.0	1	3	1	41	1	-1	0	0
29999	30000	50000.0	1	2	1	46	0	0	0	0

30000 rows × 25 columns

Printing the Dataset Information

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 25 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   ID                                         30000 non-null  int64
1   LIMIT_BAL                                30000 non-null  float64
2   SEX                                       30000 non-null  int64
3   EDUCATION                                30000 non-null  int64
4   MARRIAGE                                 30000 non-null  int64
5   AGE                                       30000 non-null  int64
6   PAY_0                                    30000 non-null  int64
7   PAY_2                                    30000 non-null  int64
8   PAY_3                                    30000 non-null  int64
9   PAY_4                                    30000 non-null  int64
10  PAY_5                                    30000 non-null  int64
11  PAY_6                                    30000 non-null  int64
12  BILL_AMT1                               30000 non-null  float64
13  BILL_AMT2                               30000 non-null  float64
14  BILL_AMT3                               30000 non-null  float64
15  BILL_AMT4                               30000 non-null  float64
16  BILL_AMT5                               30000 non-null  float64
17  BILL_AMT6                               30000 non-null  float64
18  PAY_AMT1                                30000 non-null  float64
19  PAY_AMT2                                30000 non-null  float64
20  PAY_AMT3                                30000 non-null  float64
21  PAY_AMT4                                30000 non-null  float64
22  PAY_AMT5                                30000 non-null  float64
23  PAY_AMT6                                30000 non-null  float64
24  default.payment.next.month              30000 non-null  int64
dtypes: float64(13), int64(12)
memory usage: 5.7 MB
```

In our dataset we got customer credit card transaction history for past 6 month , on basis of which we have to predict if customer will default or not.

Checking for any null values in the dataset

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: ID
LIMIT_BAL
SEX
EDUCATION
MARRIAGE
AGE
PAY_0
PAY_2
PAY_3
PAY_4
PAY_5
PAY_6
BILL_AMT1
BILL_AMT2
BILL_AMT3
BILL_AMT4
BILL_AMT5
BILL_AMT6
PAY_AMT1
PAY_AMT2
PAY_AMT3
PAY_AMT4
PAY_AMT5
PAY_AMT6
default.payment.next.month
dtype: int64
```

Data Description


Here we check on data description to the descriptive features like mean,max,std and others from our dataset

```
In [ ]: df.describe()
```

Out[]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	15000.500000	167484.322667	1.603733	1.853133	1.551867	35.485500
std	8660.398374	129747.661567	0.489129	0.790349	0.521970	9.217904
min	1.000000	10000.000000	1.000000	0.000000	0.000000	21.000000
25%	7500.750000	50000.000000	1.000000	1.000000	1.000000	28.000000
50%	15000.500000	140000.000000	2.000000	2.000000	2.000000	34.000000
75%	22500.250000	240000.000000	2.000000	2.000000	2.000000	41.000000
max	30000.000000	1000000.000000	2.000000	6.000000	3.000000	79.000000

8 rows × 25 columns



Exploratory Data Analysis

Dependent Variable:

```
In [ ]: #renaming for better convinience
df['IsDefaulter'] =df ['default.payment.next.month']
df.drop('default.payment.next.month',axis = 1)
# df.rename({'default.payment.next.month' : 'IsDefaulter'}, inplace=True)
```


```
Out[ ]:
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4
0	1	20000.0	2	2	1	24	2	2	-1	-1
1	2	120000.0	2	2	2	26	-1	2	0	0
2	3	90000.0	2	2	2	34	0	0	0	0
3	4	50000.0	2	2	1	37	0	0	0	0
4	5	50000.0	1	2	1	57	-1	0	-1	0
...
29995	29996	220000.0	1	3	1	39	0	0	0	0
29996	29997	150000.0	1	3	2	43	-1	-1	-1	-1
29997	29998	30000.0	1	2	2	37	4	3	2	-1
29998	29999	80000.0	1	3	1	41	1	-1	0	0
29999	30000	50000.0	1	2	1	46	0	0	0	0

30000 rows × 25 columns

```
In [ ]: plt.figure(figsize=(10,5))
sns.countplot(x = 'IsDefaulter', data = df)
```

```
Out[ ]: <Axes: xlabel='IsDefaulter', ylabel='count'>
```



```
In [ ]: df['IsDefaulter'].value_counts()
```

```
Out[ ]: IsDefaulter
0      23364
1       6636
Name: count, dtype: int64
```

As we can see from above graph that both classes are not in proportion and we have imbalanced dataset.

Independent Variable:

Categorical Features

We have few categorical features in our dataset.They are demonstrated as below.

SEX

- 1 - Male
- 2 - Female

```
In [ ]: df['SEX'].value_counts()
```

```
Out[ ]: SEX
2      18112
1      11888
Name: count, dtype: int64
```

Education

1 = graduate school; 2 = university; 3 = high school; 4 = others

```
In [ ]: df['EDUCATION'].value_counts()
```

```
Out[ ]: EDUCATION
2      14030
1      10585
3       4917
4        468
Name: count, dtype: int64
```

As we can see in dataset we have values like 5,6,0 as well for which we are not having description so we can add up them in 4, which is Others.

```
In [ ]: fil = (df['EDUCATION'] == 5) | (df['EDUCATION'] == 6) | (df['EDUCATION'] == 0)
df.loc[fil, 'EDUCATION'] = 4
df['EDUCATION'].value_counts()
```

```
Out[ ]: EDUCATION
2      14030
1      10585
3       4917
4        468
Name: count, dtype: int64
```

Marriage

1 = married; 2 = single; 3 = others

```
In [ ]: df['MARRIAGE'].value_counts()
```

```
Out[ ]: MARRIAGE
2      15964
1      13659
3        377
Name: count, dtype: int64
```

We have few values for 0, which are not determined therefore the are added in Others category.

```
In [ ]: fil = df['MARRIAGE'] == 0
df.loc[fil, 'MARRIAGE'] = 3
df['MARRIAGE'].value_counts()
```

```
Out[ ]: MARRIAGE
2      15964
1      13659
3         377
Name: count, dtype: int64
```

Plotting our categorical features


```
In [ ]: categorical_features = ['SEX', 'EDUCATION', 'MARRIAGE']
```

```
In [ ]: df_cat = df[categorical_features]
df_cat['Defaulter'] = df['IsDefaulter']
```


```
In [ ]: df_cat.replace({'SEX': {1 : 'MALE', 2 : 'FEMALE'}}, 'EDUCATION' : {1 : 'grad
```

```
In [ ]: for col in categorical_features:
    plt.figure(figsize=(10,5))
    fig, axes = plt.subplots(ncols=2,figsize=(13,8))
    df[col].value_counts().plot(kind="pie",ax = axes[0],subplots=True)
    sns.countplot(x = col, hue = 'Defaulter', data = df_cat)
```


<Figure size 1000x500 with 0 Axes>



<Figure size 1000x500 with 0 Axes>



<Figure size 1000x500 with 0 Axes>



Below are few observations for categorical features:

- There are more females credit card holder,so no. of defaulter have high proportion of females.
- No. of defaulters have a higher proportion of educated people (graduate school and university)
- No. of defaulters have a higher proportion of Singles.

Limit Balance

```
In [ ]: df['LIMIT_BAL'].max()
```

```
Out[ ]: 1000000.0
```

```
In [ ]: df['LIMIT_BAL'].min()
```


```
Out[ ]: 10000.0
```

```
In [ ]: df['LIMIT_BAL'].describe()
```


```
Out[ ]: count      30000.000000  
mean      167484.322667  
std      129747.661567  
min       10000.000000  
25%       50000.000000  
50%      140000.000000  
75%      240000.000000  
max      1000000.000000  
Name: LIMIT_BAL, dtype: float64
```

```
In [ ]: sns.barplot(x='IsDefaulter', y='LIMIT_BAL', data=df)
```

```
Out[ ]: <Axes: xlabel='IsDefaulter', ylabel='LIMIT_BAL'>
```



```
In [ ]: plt.figure(figsize=(10,10))  
ax = sns.boxplot(x="IsDefaulter", y="LIMIT_BAL", data=df)
```

```
In [ ]: #renaming columns


df.rename(columns={'PAY_0': 'PAY_SEPT', 'PAY_2': 'PAY_AUG', 'PAY_3': 'PAY_JUL', '
df.rename(columns={'BILL_AMT1': 'BILL_AMT_SEPT', 'BILL_AMT2': 'BILL_AMT_AUG', '
df.rename(columns={'PAY_AMT1': 'PAY_AMT_SEPT', 'PAY_AMT2': 'PAY_AMT_AUG', 'PAY_
```

```
In [ ]: df.head()
```

Out[]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_SEPT	PAY_AUG	PAY_JUL	PAY_J
0	1	20000.0	2	2	1	24	2	2	-1	
1	2	120000.0	2	2	2	26	-1	2	0	
2	3	90000.0	2	2	2	34	0	0	0	
3	4	50000.0	2	2	1	37	0	0	0	
4	5	50000.0	1	2	1	57	-1	0	-1	

5 rows × 26 columns



AGE

Plotting graph of number of ages of all people with credit card irrespective of gender.


```
In [ ]: df['AGE'].value_counts()
```

```
Out[ ]: AGE
29      1605
27      1477
28      1409
30      1395
26      1256
31      1217
25      1186
34      1162
32      1158
33      1146
24      1127
35      1113
36      1108
37      1041
39       954
38       944
23       931
40       870
41       824
42       794
44       700
43       670
45       617
46       570
22       560
47       501
48       466
49       452
50       411
51       340
53       325
52       304
54       247
55       209
56       178
58       122
57       122
59        83
60        67
21        67
61        56
62        44
63        31
64        31
66        25
65        24
67        16
69        15
70        10
68         5
73         4
72         3
75         3
71         3
79         1
74         1
Name: count, dtype: int64
```

```
In [ ]: df['AGE']=df['AGE'].astype('int')
```

```
In [ ]: # subplots of age
plt.figure(figsize=(20,10))
sns.countplot(x='AGE',hue='IsDefaulter',data=df)
```

```
Out[ ]: <Axes: xlabel='AGE', ylabel='count'>
```




```
In [ ]: df.groupby('IsDefaulter')['AGE'].mean()
```

```
Out[ ]: IsDefaulter
0      35.417266
1      35.725738
Name: AGE, dtype: float64
```

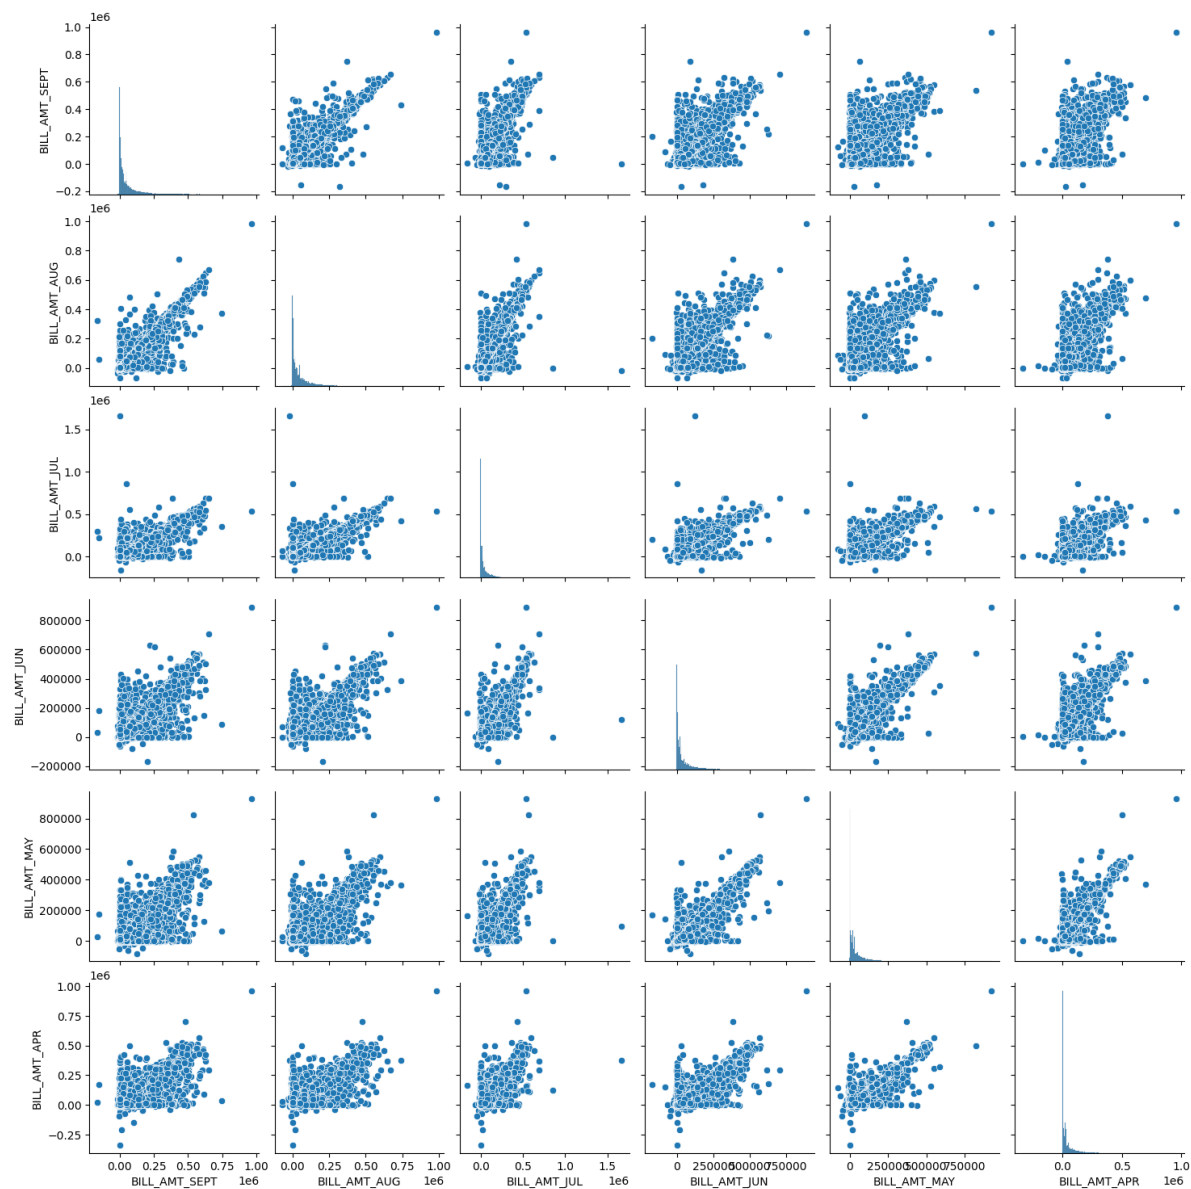
```
In [ ]: df = df.astype('int')
```

```
In [ ]: plt.figure(figsize=(10,10))
ax = sns.boxplot(x="IsDefaulter", y="AGE", data=df)
```




Bill Amount




```
In [ ]: bill_amnt_df = df[['BILL_AMT_SEPT', 'BILL_AMT_AUG', 'BILL_AMT_JUL', 'BI  
In [ ]: sns.pairplot(data = bill_amnt_df)  
Out[ ]: <seaborn.axisgrid.PairGrid at 0x7f1bbb6a34f0>
```





History payment status

```
In [ ]: pay_col = ['PAY_SEPT', 'PAY_AUG', 'PAY_JUL', 'PAY_JUN', 'PA
for col in pay_col:
    plt.figure(figsize=(10,5))
    sns.countplot(x = col, hue = 'IsDefaulter', data = df)
```








Paid Amount

```
In [ ]: pay_amnt_df = df[['PAY_AMT_SEPT', 'PAY_AMT_AUG', 'PAY_AMT_JUL', 'PA
```

```
In [ ]: sns.pairplot(data = pay_amnt_df, hue='IsDefaulter')
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x7f2cd5589d20>
```



```
In [ ]: # shape of the data
df.shape
```

```
Out[ ]: (30000, 26)
```

As we have seen earlier that we have imbalanced dataset. So to remediate Imbalance we are using SMOTE(Synthetic Minority Oversampling Technique)

```
In [ ]: from imblearn.over_sampling import SMOTE

smote = SMOTE()

# fit predictor and target variable
x_smote, y_smote = smote.fit_resample(df.iloc[:,0:-1], df['IsDefaulter'])

print('Original dataset shape', len(df))
print('Resampled dataset shape', len(y_smote))
```

```
Original dataset shape 30000
Resampled dataset shape 46728
```

```
In [ ]: x_smote
```



```
Out[ ]:
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_SEPT	PAY_AUG	PAY_JUL
0	1	20000	2	2	1	24	2	2	-1
1	2	120000	2	2	2	26	-1	2	0
2	3	90000	2	2	2	34	0	0	0
3	4	50000	2	2	1	37	0	0	0
4	5	50000	1	2	1	57	-1	0	-1
...
46723	10827	80000	2	2	1	29	0	0	-1
46724	3883	20000	1	1	1	31	2	2	2
46725	6998	130000	1	1	2	34	0	0	0
46726	25713	30000	2	2	1	37	2	2	2
46727	16177	383043	2	1	1	36	-1	-1	-1

46728 rows × 25 columns

```
In [ ]: columns = list(df.columns)
```

```
In [ ]: columns.pop()
```

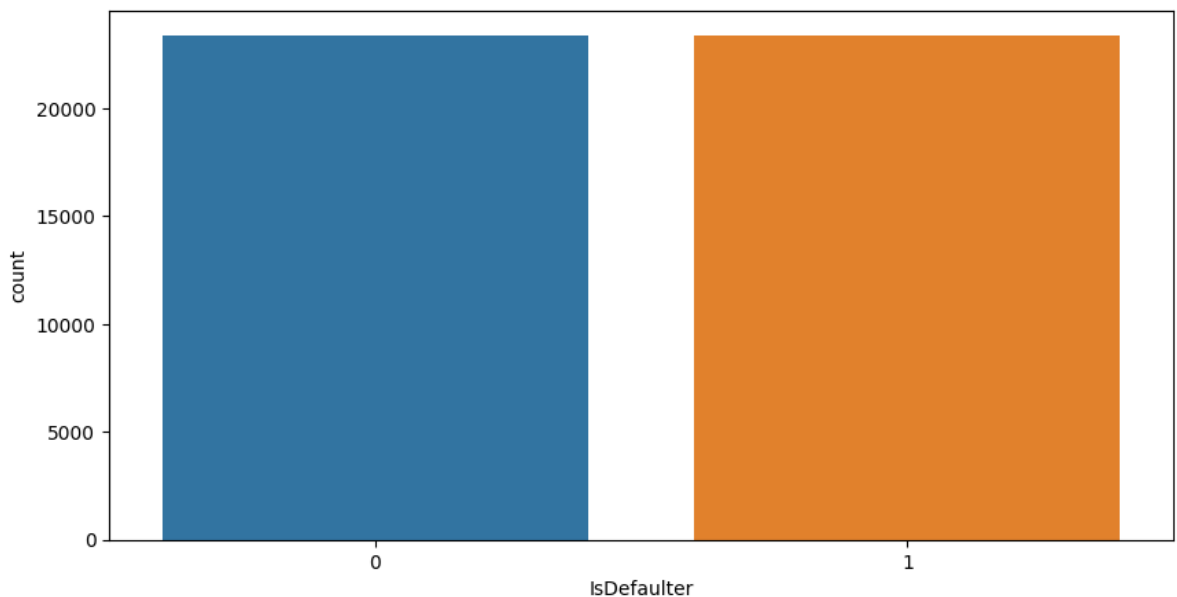
```
Out[ ]: 'IsDefaulter'
```

```
In [ ]: balance_df = pd.DataFrame(x_smote, columns=columns)
```

```
In [ ]: balance_df['IsDefaulter'] = y_smote
```

```
In [ ]: # plot the count after resample
plt.figure(figsize=(10,5))
sns.countplot(x = 'IsDefaulter', data = balance_df)
```

```
Out[ ]: <Axes: xlabel='IsDefaulter', ylabel='count'>
```



```
In [ ]: balance_df[balance_df['IsDefaulter']==1]
```

Out[]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_SEPT	PAY_AUG	PAY_JUL
0	1	20000	2	2	1	24	2	2	-1
1	2	120000	2	2	2	26	-1	2	0
13	14	70000	1	2	2	30	1	2	2
16	17	20000	1	1	2	24	0	0	2
21	22	120000	2	2	1	39	-1	-1	-1
...
46723	19742	56575	2	2	2	26	1	0	0
46724	7160	252470	1	2	1	53	0	0	0
46725	23249	20000	2	2	2	44	0	0	0
46726	29479	10000	1	2	2	33	0	0	0
46727	26037	30000	1	2	1	37	1	1	1

23364 rows × 26 columns

Feature Engineering

```

In [ ]: df_fr = balance_df.copy()

In [ ]: df_fr['Payment_Value'] = df_fr['PAY_SEPT'] + df_fr['PAY_AUG'] + df_fr['PAY_JUL']


In [ ]: df_fr.groupby('IsDefaulter')['Payment_Value'].mean()

Out[ ]: IsDefaulter
0    -1.980140
1     1.704503
Name: Payment_Value, dtype: float64

In [ ]: plt.figure(figsize=(10,10))
sns.boxplot(data = df_fr, x = 'IsDefaulter', y = 'Payment_Value' )

Out[ ]: <Axes: xlabel='IsDefaulter', ylabel='Payment_Value'>

```



```
In [ ]: df_fr['Dues'] = (df_fr['BILL_AMT_APR']+df_fr['BILL_AMT_MAY']+df_fr['BILL_AMT_JUN']
```

```
In [ ]: df_fr.groupby('IsDefaulter')['Dues'].mean()
```

```
Out[ ]: IsDefaulter
0      187742.051532
1      195826.211822
Name: Dues, dtype: float64
```

```
In [ ]: df_fr['EDUCATION'].unique()
```

```
Out[ ]: array([2, 1, 3, 4])
```

```
In [ ]: df_fr['EDUCATION']=np.where(df_fr['EDUCATION'] == 6, 4, df_fr['EDUCATION'])
df_fr['EDUCATION']=np.where(df_fr['EDUCATION'] == 0, 4, df_fr['EDUCATION'])
```

```
In [ ]: df_fr['MARRIAGE'].unique()
```

```
Out[ ]: array([1, 2, 3])
```

```
In [ ]: df_fr['MARRIAGE']=np.where(df_fr['MARRIAGE'] == 0, 3, df_fr['MARRIAGE'])
```

```
In [ ]: df_fr.replace({'SEX': {1 : 'MALE', 2 : 'FEMALE'}, 'EDUCATION' : {1 : 'gradu
```

```
In [ ]: df_fr.head()
```

Out[]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_SEPT	PAY_AUG	PAY_JUL	PAY_JAN
0	1	20000	FEMALE	university	married	24	2	2	-1	
1	2	120000	FEMALE	university	single	26	-1	2	0	
2	3	90000	FEMALE	university	single	34	0	0	0	
3	4	50000	FEMALE	university	married	37	0	0	0	
4	5	50000	MALE	university	married	57	-1	0	-1	

5 rows × 28 columns

One Hot Encoding

```
In [ ]: df_fr = pd.get_dummies(df_fr,columns=['EDUCATION','MARRIAGE'])
```

```
In [ ]: df_fr.head()
```

Out[]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_SEPT	PAY_AUG	PAY_JUL	PAY_JAN
0	1	20000	FEMALE	university	married	24	2	2	-1	
1	2	120000	FEMALE	university	single	26	-1	2	0	
2	3	90000	FEMALE	university	single	34	0	0	0	
3	4	50000	FEMALE	university	married	37	0	0	0	
4	5	50000	MALE	university	married	57	-1	0	-1	

5 rows × 28 columns

```
In [ ]: df_fr.drop(['EDUCATION_others','MARRIAGE_others'],axis = 1, inplace = True)
df_fr = pd.get_dummies(df_fr, columns = ['PAY_SEPT','PAY_AUG','PAY_JUL','PAY_JAN'])
df_fr.head()
```

Out[]:

	ID	LIMIT_BAL	SEX	AGE	BILL_AMT_SEPT	BILL_AMT_AUG	BILL_AMT_JUL	BILL_AMT_JAN
0	1	20000	FEMALE	24	3913	3102	689	
1	2	120000	FEMALE	26	2682	1725	2682	
2	3	90000	FEMALE	34	29239	14027	13559	
3	4	50000	FEMALE	37	46990	48233	49291	
4	5	50000	MALE	57	8617	5670	35835	

5 rows × 85 columns

```
In [ ]: # LABEL ENCODING FOR SEX
encoders_nums = {
    "SEX":{"FEMALE": 0, "MALE": 1}
}
df_fr = df_fr.replace(encoders_nums)
df_fr.head()
```

```
Out[ ]:
```

	ID	LIMIT_BAL	SEX	AGE	BILL_AMT_SEPT	BILL_AMT_AUG	BILL_AMT_JUL	BILL_AMT_JUL
0	1	20000	0	24	3913	3102	689	
1	2	120000	0	26	2682	1725	2682	327
2	3	90000	0	34	29239	14027	13559	1433
3	4	50000	0	37	46990	48233	49291	2831
4	5	50000	1	57	8617	5670	35835	2094

5 rows × 85 columns

```
In [ ]: df_fr.drop('ID',axis = 1, inplace = True)
df_fr.to_csv('Final_df.csv')
df_fr = pd.read_csv('./Final_df.csv')

df_fr.head()
```

```
Out[ ]:
```

	Unnamed: 0	LIMIT_BAL	SEX	AGE	BILL_AMT_SEPT	BILL_AMT_AUG	BILL_AMT_JUL	BILL_
0	0	20000	0	24	3913	3102	689	
1	1	120000	0	26	2682	1725	2682	
2	2	90000	0	34	29239	14027	13559	
3	3	50000	0	37	46990	48233	49291	
4	4	50000	1	57	8617	5670	35835	

5 rows × 85 columns

```
In [ ]: df_fr.drop(['Unnamed: 0'],axis = 1, inplace = True)
```

Implementing Logistic Regression

```
In [ ]: df_log_reg = df_fr.copy()
```

```
In [ ]: df_log_reg.head()
```

```
Out[ ]:
```

	LIMIT_BAL	SEX	AGE	BILL_AMT_SEPT	BILL_AMT_AUG	BILL_AMT_JUL	BILL_AMT_JUN	E
0	20000	0	24	3913	3102	689		0
1	120000	0	26	2682	1725	2682		3272
2	90000	0	34	29239	14027	13559		14331
3	50000	0	37	46990	48233	49291		28314
4	50000	1	57	8617	5670	35835		20940

5 rows × 84 columns

```
In [ ]: X = df_log_reg.drop(['IsDefaulter','Payment_Value','Dues'],axis=1)
y = df_log_reg['IsDefaulter']
```

```
columns = X.columns  
scaler = StandardScaler()  
X = scaler.fit_transform(X)
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, r
```

```
In [ ]: param_grid = {'penalty':['l1','l2'], 'C' : [0.001, 0.01, 0.1, 1, 10, 100, 1
```

```
In [ ]: grid_lr_clf = GridSearchCV(LogisticRegression(), param_grid, scoring = 'acc  
grid_lr_clf.fit(X_train, y_train)
```

Fitting 3 folds for each of 14 candidates, totalling 42 fits

```
[CV 3/3] END .....C=0.001, penalty=l1;, score=nan total time=
0.1s
[CV 1/3] END .....C=0.001, penalty=l1;, score=nan total time=
0.1s
[CV 2/3] END .....C=0.001, penalty=l1;, score=nan total time=
0.1s
[CV 1/3] END .....C=0.01, penalty=l1;, score=nan total time=
0.0s
[CV 2/3] END .....C=0.01, penalty=l1;, score=nan total time=
0.0s
[CV 3/3] END .....C=0.01, penalty=l1;, score=nan total time=
0.0s
[CV 2/3] END .....C=0.001, penalty=l2;, score=1.000 total time=
5.2s
[CV 3/3] END .....C=0.001, penalty=l2;, score=1.000 total time=
5.1s
[CV 1/3] END .....C=0.01, penalty=l2;, score=1.000 total time=
5.0s
[CV 1/3] END .....C=0.001, penalty=l2;, score=1.000 total time=
5.4s
[CV 1/3] END .....C=0.1, penalty=l1;, score=nan total time=
0.1s
[CV 2/3] END .....C=0.1, penalty=l1;, score=nan total time=
0.1s
[CV 3/3] END .....C=0.1, penalty=l1;, score=nan total time=
0.1s
[CV 3/3] END .....C=0.01, penalty=l2;, score=1.000 total time=
0.4s
[CV 2/3] END .....C=0.01, penalty=l2;, score=1.000 total time=
0.5s
[CV 1/3] END .....C=1, penalty=l1;, score=nan total time=
0.0s
[CV 2/3] END .....C=1, penalty=l1;, score=nan total time=
0.0s
[CV 3/3] END .....C=1, penalty=l1;, score=nan total time=
0.0s
[CV 1/3] END .....C=0.1, penalty=l2;, score=1.000 total time=
0.4s
[CV 3/3] END .....C=0.1, penalty=l2;, score=1.000 total time=
0.4s
[CV 2/3] END .....C=0.1, penalty=l2;, score=1.000 total time=
0.7s
[CV 1/3] END .....C=10, penalty=l1;, score=nan total time=
0.1s
[CV 2/3] END .....C=10, penalty=l1;, score=nan total time=
0.0s
[CV 3/3] END .....C=10, penalty=l1;, score=nan total time=
0.0s
[CV 1/3] END .....C=1, penalty=l2;, score=1.000 total time=
0.6s
[CV 2/3] END .....C=1, penalty=l2;, score=1.000 total time=
0.7s
[CV 3/3] END .....C=1, penalty=l2;, score=1.000 total time=
0.7s
[CV 1/3] END .....C=100, penalty=l1;, score=nan total time=
0.1s
[CV 2/3] END .....C=100, penalty=l1;, score=nan total time=
0.0s
[CV 3/3] END .....C=100, penalty=l1;, score=nan total time=
0.0s
[CV 1/3] END .....C=10, penalty=l2;, score=1.000 total time=
0.6s
[CV 3/3] END .....C=10, penalty=l2;, score=1.000 total time=
```

```

0.6s
[CV 2/3] END .....C=10, penalty=l2;, score=1.000 total time=
0.9s
[CV 1/3] END .....C=100, penalty=l2;, score=1.000 total time=
0.5s
[CV 2/3] END .....C=100, penalty=l2;, score=1.000 total time=
0.5s
[CV 3/3] END .....C=1000, penalty=l1;, score=nan total time=
0.1s
[CV 2/3] END .....C=1000, penalty=l1;, score=nan total time=
0.2s
[CV 3/3] END .....C=100, penalty=l2;, score=1.000 total time=
0.7s
[CV 1/3] END .....C=1000, penalty=l2;, score=1.000 total time=
0.4s
[CV 1/3] END .....C=1000, penalty=l1;, score=nan total time=
0.8s
[CV 2/3] END .....C=1000, penalty=l2;, score=1.000 total time=
0.4s
[CV 3/3] END .....C=1000, penalty=l2;, score=1.000 total time=
0.3s

```

```

Out[ ]:
└─ GridSearchCV
  └─ estimator: LogisticRegression
    └─ LogisticRegression

```

```
In [ ]: optimized_clf = grid_lr_clf.best_estimator_
```

```
In [ ]: grid_lr_clf.best_params_
```

```
Out[ ]: {'C': 0.01, 'penalty': 'l2'}
```

```
In [ ]: grid_lr_clf.best_score_
```

```
Out[ ]: 1.0
```

```

In [ ]: # Predicted Probability
train_preds = optimized_clf.predict_proba(X_train)[:,-1]
test_preds = optimized_clf.predict_proba(X_test)[:,-1]
# Get the predicted classes
train_class_preds = optimized_clf.predict(X_train)
test_class_preds = optimized_clf.predict(X_test)
# Get the accuracy scores
train_accuracy_lr = accuracy_score(train_class_preds,y_train)
test_accuracy_lr = accuracy_score(test_class_preds,y_test)

print("The accuracy on train data is ", train_accuracy_lr)
print("The accuracy on test data is ", test_accuracy_lr)

```

```

The accuracy on train data is  1.0
The accuracy on test data is  1.0

```

```

In [ ]: test_accuracy_lr = accuracy_score(test_class_preds,y_test)
test_precision_score_lr = precision_score(test_class_preds,y_test)
test_recall_score_lr = recall_score(test_class_preds,y_test)
test_f1_score_lr = f1_score(test_class_preds,y_test)
test_roc_score_lr = roc_auc_score(test_class_preds,y_test)

print("The accuracy on test data is ", test_accuracy_lr)
print("The precision on test data is ", test_precision_score_lr)

```



```
print("The recall on test data is ", test_recall_score_lr)
print("The f1 on test data is ", test_f1_score_lr)
print("The roc_score on test data is ", test_roc_score_lr)
```

```
The accuracy on test data is  1.0
The precision on test data is  1.0
The recall on test data is  1.0
The f1 on test data is  1.0
The roc_score on test data is  1.0
```

```
In [ ]: # Get the confusion matrix for both train and test
```


```
labels = ['Not Defaulter', 'Defaulter']
cm = confusion_matrix(y_train, train_class_preds)
print(cm)

ax= plt.subplot()
sns.heatmap(cm, annot=True, ax = ax) #annot=True to annotate cells

# labels, title and ticks
ax.set_xlabel('Predicted labels')
ax.set_ylabel('True labels')
ax.set_title('Confusion Matrix')
ax.xaxis.set_ticklabels(labels)
ax.yaxis.set_ticklabels(labels)
```

```
[[15653    0]
 [    0 15654]]
```


```
Out[ ]: [Text(0, 0.5, 'Not Defaulter'), Text(0, 1.5, 'Defaulter')]
```



```
In [ ]: feature_importance = pd.DataFrame({'Features':columns, 'Importance':np.abs(
```


```
In [ ]: feature_importance = feature_importance.sort_values(by = 'Importance', asce
```

```
In [ ]: plt.bar(height=feature_importance['Importance'], x= feature_importance['Fea
plt.xticks(rotation=80)
plt.title("Feature importances via coefficients")
plt.show()
```



```
In [ ]: y_preds_proba_lr = optimized_clf.predict_proba(X_test)[::,1]
```

```
In [ ]: y_pred_proba = y_preds_proba_lr
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
auc = roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



We have implemented logistic regression and we getting f1-score approx 73%. As we have imbalanced dataset, F1- score is better parameter.

Implementing SVC

```
In [ ]: from sklearn.model_selection import GridSearchCV
```

```
In [ ]: # defining parameter range
param_grid = {'C': [0.1, 1, 10, 100],
              'kernel': ['rbf']}
X = df_fr.drop(['IsDefaulter', 'Payement_Value', 'Dues'], axis=1)
y = df_fr['IsDefaulter']
scaler = StandardScaler()
X = scaler.fit_transform(X)
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, r
```

```
In [ ]: grid_clf = GridSearchCV(SVC(probability=True), param_grid, scoring = 'accu
grid_clf.fit(X_train, y_train)
```

Fitting 3 folds for each of 4 candidates, totalling 12 fits

```
[CV 1/3] END .....C=1, kernel=rbf;, score=0.993 total time=
1.3min
[CV 2/3] END .....C=1, kernel=rbf;, score=0.995 total time=
1.4min
[CV 3/3] END .....C=0.1, kernel=rbf;, score=0.988 total time=
3.0min
[CV 2/3] END .....C=0.1, kernel=rbf;, score=0.987 total time=
3.2min
[CV 1/3] END .....C=0.1, kernel=rbf;, score=0.985 total time=
3.4min
[CV 1/3] END .....C=10, kernel=rbf;, score=0.995 total time=
1.0min
[CV 3/3] END .....C=1, kernel=rbf;, score=0.995 total time=
1.4min
[CV 2/3] END .....C=10, kernel=rbf;, score=0.997 total time=
1.1min
[CV 3/3] END .....C=10, kernel=rbf;, score=0.997 total time=
1.1min
[CV 1/3] END .....C=100, kernel=rbf;, score=0.995 total time=
48.0s
[CV 2/3] END .....C=100, kernel=rbf;, score=0.997 total time=
50.6s
[CV 3/3] END .....C=100, kernel=rbf;, score=0.997 total time=
47.7s
```

```
Out[ ]:  ▶ GridSearchCV
          ▶ estimator: SVC
            ▶ SVC
```

```
In [ ]: optimal_SVC_clf = grid_clf.best_estimator_
```

```
In [ ]: grid_clf.best_params_
```

```
Out[ ]: {'C': 100, 'kernel': 'rbf'}
```

```
In [ ]: grid_clf.best_score_
```

```
Out[ ]: 0.9964544823494704
```

```
In [ ]: # Get the predicted classes
train_class_preds = optimal_SVC_clf.predict(X_train)
test_class_preds = optimal_SVC_clf.predict(X_test)
```

```
In [ ]: # Get the accuracy scores
train_accuracy_SVC = accuracy_score(train_class_preds,y_train)
test_accuracy_SVC = accuracy_score(test_class_preds,y_test)

print("The accuracy on train data is ", train_accuracy_lr)
print("The accuracy on test data is ", test_accuracy_lr)
```

```
The accuracy on train data is  1.0
The accuracy on test data is  1.0
```

```
In [ ]: test_accuracy_SVC = accuracy_score(test_class_preds,y_test)
test_precision_score_SVC = precision_score(test_class_preds,y_test)
test_recall_score_SVC = recall_score(test_class_preds,y_test)
test_f1_score_SVC = f1_score(test_class_preds,y_test)
test_roc_score_SVC = roc_auc_score(test_class_preds,y_test)
```

```
print("The accuracy on test data is ", test_accuracy_SVC)
print("The precision on test data is ", test_precision_score_SVC)
print("The recall on test data is ", test_recall_score_SVC)
print("The f1 on test data is ", test_f1_score_SVC)
print("The roc_score on test data is ", test_roc_score_SVC)
```

```
The accuracy on test data is  0.9968873613903119
The precision on test data is  0.9990920881971466
The recall on test data is  0.9947055785123967
The f1 on test data is  0.9968940080238126
The roc_score on test data is  0.9968968820007601
```

We can see from above results that we are getting around 80% train accuracy and 78% for test accuracy which is not bad. But f1- score is 76% approx, so there might be more ground for improvement.


```
In [ ]:  # Get the confusion matrix for both train and test

labels = ['Not Defaulter', 'Defaulter']
cm = confusion_matrix(y_train, train_class_preds)
print(cm)

ax= plt.subplot()
sns.heatmap(cm, annot=True, ax = ax) #annot=True to annotate cells

# labels, title and ticks
ax.set_xlabel('Predicted labels')
ax.set_ylabel('True labels')
ax.set_title('Confusion Matrix')
ax.xaxis.set_ticklabels(labels)
ax.yaxis.set_ticklabels(labels)

[[15653      0]
 [      0 15654]]
Out[ ]: [Text(0, 0.5, 'Not Defaulter'), Text(0, 1.5, 'Defaulter')]
```



```
In [ ]: import torch
```

```
In [ ]: model_save_name = 'SVC_optimized_classifier.pt'
path = F"./{model_save_name}"
torch.save(optimal_SVC_clf, path)
```

```
In [ ]: model_save_name = 'SVC_optimized_classifier.pt'
path = F"./{model_save_name}"
optimal_SVC_clf = torch.load(path)
```


```
In [ ]: optimal_SVC_clf
```

```
Out[ ]: SVC
SVC(C=100, probability=True)
```

```
In [ ]: # Get the predicted classes
train_class_preds = optimal_SVC_clf.predict(X_train)
test_class_preds = optimal_SVC_clf.predict(X_test)
```

```
In [ ]: y_pred_proba_SVC = optimal_SVC_clf.predict_proba(X_test)[::,1]
```

```
In [ ]: # ROC AUC CURVE
fpr, tpr, _ = roc_curve(y_test, y_pred_proba_SVC)
auc = roc_auc_score(y_test, y_pred_proba_SVC)
plt.plot(fpr, tpr, label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



Implementing Decision Tree

Decision Tree is another very popular algorithm for classification problems because it is easy to interpret and understand. An internal node represents a feature, the branch represents a decision rule, and each leaf node represents the outcome. Some advantages of decision trees are that they require less data preprocessing, i.e., no need to normalize features. However, noisy data can be easily overfitted and results in biased results when the data set is imbalanced.

```
In [ ]: param_grid = {'max_depth': [20,30,50,100], 'min_samples_split':[0.1,0.2,0.4
```

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
```

```
In [ ]: X = df_fr.drop(['IsDefaulter', 'Payement_Value', 'Dues'],axis=1)
        y = df_fr['IsDefaulter']
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, r
```

```
In [ ]: grid_DTC_clf = GridSearchCV(DecisionTreeClassifier(), param_grid, scoring =
        grid_DTC_clf.fit(X_train, y_train)
```

Fitting 3 folds for each of 12 candidates, totalling 36 fits

```
[CV 1/3] END max_depth=20, min_samples_split=0.1, score=1.000 total time=0.2s
[CV 2/3] END max_depth=20, min_samples_split=0.1, score=1.000 total time=0.2s
[CV 3/3] END max_depth=20, min_samples_split=0.1, score=1.000 total time=0.2s
[CV 2/3] END max_depth=20, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 1/3] END max_depth=20, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 3/3] END max_depth=20, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 2/3] END max_depth=20, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 1/3] END max_depth=20, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 3/3] END max_depth=20, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 1/3] END max_depth=30, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 2/3] END max_depth=30, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 1/3] END max_depth=30, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 3/3] END max_depth=30, min_samples_split=0.1, score=1.000 total time=0.2s
[CV 2/3] END max_depth=30, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 1/3] END max_depth=30, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 3/3] END max_depth=30, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 3/3] END max_depth=30, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 2/3] END max_depth=30, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 1/3] END max_depth=50, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 2/3] END max_depth=50, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 3/3] END max_depth=50, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 1/3] END max_depth=50, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 2/3] END max_depth=50, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 3/3] END max_depth=50, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 1/3] END max_depth=50, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 3/3] END max_depth=50, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 2/3] END max_depth=50, min_samples_split=0.4, score=1.000 total time=0.1s
[CV 1/3] END max_depth=100, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 2/3] END max_depth=100, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 3/3] END max_depth=100, min_samples_split=0.1, score=1.000 total time=0.1s
[CV 2/3] END max_depth=100, min_samples_split=0.2, score=1.000 total time=0.1s
[CV 1/3] END max_depth=100, min_samples_split=0.2, score=1.000 total time=
```



```

0.2s
[CV 3/3] END max_depth=100, min_samples_split=0.2;; score=1.000 total time=
0.1s
[CV 2/3] END max_depth=100, min_samples_split=0.4;; score=1.000 total time=
0.1s
[CV 1/3] END max_depth=100, min_samples_split=0.4;; score=1.000 total time=
0.1s[CV 3/3] END max_depth=100, min_samples_split=0.4;; score=1.000 total t
ime= 0.1s

```

```

Out[ ]:
└─ GridSearchCV
  └─ estimator: DecisionTreeClassifier
    └─ DecisionTreeClassifier

```

```
In [ ]: grid_DTC_clf.best_score_
```

```
Out[ ]: 1.0
```

```
In [ ]: optimal_DTC_clf = grid_DTC_clf.best_estimator_
```

```
In [ ]: # Get the predicted classes
train_class_preds = optimal_DTC_clf.predict(X_train)
test_class_preds = optimal_DTC_clf.predict(X_test)
```

```
In [ ]: grid_DTC_clf.best_params_
```

```
Out[ ]: {'max_depth': 20, 'min_samples_split': 0.1}
```

```
In [ ]: # Get the accuracy scores
train_accuracy_DTC = accuracy_score(train_class_preds,y_train)
test_accuracy_DTC = accuracy_score(test_class_preds,y_test)

print("The accuracy on train data is ", train_accuracy_DTC)
print("The accuracy on test data is ", test_accuracy_DTC)
```

```

The accuracy on train data is  1.0
The accuracy on test data is  1.0

```

Implementing RandomForest

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
```

```
In [ ]: X = df_fr.drop(['IsDefaulter', 'Payement_Value', 'Dues'],axis=1)
y = df_fr['IsDefaulter']
```

```
In [ ]: rf_clf = RandomForestClassifier()
rf_clf.fit(X_train,y_train)
```

```

Out[ ]:
└─ RandomForestClassifier
  RandomForestClassifier()

```

```
In [ ]: # Get the predicted classes
train_class_preds = rf_clf.predict(X_train)
test_class_preds = rf_clf.predict(X_test)
```

```
In [ ]: # Get the accuracy scores
train_accuracy_rf = accuracy_score(train_class_preds,y_train)
test_accuracy_rf = accuracy_score(test_class_preds,y_test)

print("The accuracy on train data is ", train_accuracy_rf)
print("The accuracy on test data is ", test_accuracy_rf)
```

The accuracy on train data is 1.0
The accuracy on test data is 0.9999351533622982

```
In [ ]: test_accuracy_rf = accuracy_score(test_class_preds,y_test)
test_precision_score_rf = precision_score(test_class_preds,y_test)
test_recall_score_rf = recall_score(test_class_preds,y_test)
test_f1_score_rf = f1_score(test_class_preds,y_test)
test_roc_score_rf = roc_auc_score(test_class_preds,y_test)

print("The accuracy on test data is ", test_accuracy_rf)
print("The precision on test data is ", test_precision_score_rf)
print("The recall on test data is ", test_recall_score_rf)
print("The f1 on test data is ", test_f1_score_rf)
print("The roc_score on test data is ", test_roc_score_rf)
```

The accuracy on test data is 0.9999351533622982
The precision on test data is 1.0
The recall on test data is 0.9998703151342239
The f1 on test data is 0.9999351533622982
The roc_score on test data is 0.9999351575671119

We can see from above results that we are getting around 99% train accuracy and 83% for test accuracy which depicts that model is overfitting. However our f1-score is around 82%, which is not bad.

```
In [ ]: param_grid = {'n_estimators': [100,150,200], 'max_depth': [10,20,30]}
```

```
In [ ]: grid_rf_clf = GridSearchCV(RandomForestClassifier(), param_grid, scoring =
grid_rf_clf.fit(X_train, y_train)
```

```

Fitting 3 folds for each of 9 candidates, totalling 27 fits
[CV 1/3] END ....max_depth=10, n_estimators=100;; score=1.000 total time=
3.8s
[CV 3/3] END ....max_depth=10, n_estimators=100;; score=1.000 total time=
4.5s
[CV 2/3] END ....max_depth=10, n_estimators=100;; score=1.000 total time=
5.0s
[CV 2/3] END ....max_depth=10, n_estimators=150;; score=1.000 total time=
7.0s
[CV 3/3] END ....max_depth=10, n_estimators=150;; score=1.000 total time=
6.8s
[CV 1/3] END ....max_depth=10, n_estimators=200;; score=1.000 total time=
10.0s
[CV 1/3] END ....max_depth=10, n_estimators=150;; score=1.000 total time=
7.3s
[CV 2/3] END ....max_depth=10, n_estimators=200;; score=1.000 total time=
9.7s
[CV 3/3] END ....max_depth=10, n_estimators=200;; score=1.000 total time=
9.8s
[CV 1/3] END ....max_depth=20, n_estimators=100;; score=1.000 total time=
6.9s
[CV 2/3] END ....max_depth=20, n_estimators=100;; score=1.000 total time=
6.1s
[CV 3/3] END ....max_depth=20, n_estimators=100;; score=1.000 total time=
5.1s
[CV 3/3] END ....max_depth=20, n_estimators=150;; score=1.000 total time=
8.4s
[CV 1/3] END ....max_depth=20, n_estimators=150;; score=1.000 total time=
10.0s
[CV 2/3] END ....max_depth=20, n_estimators=150;; score=1.000 total time=
9.5s
[CV 1/3] END ....max_depth=20, n_estimators=200;; score=1.000 total time=
9.7s
[CV 1/3] END ....max_depth=30, n_estimators=100;; score=1.000 total time=
5.6s
[CV 2/3] END ....max_depth=30, n_estimators=100;; score=1.000 total time=
5.8s
[CV 2/3] END ....max_depth=20, n_estimators=200;; score=1.000 total time=
10.3s
[CV 3/3] END ....max_depth=30, n_estimators=100;; score=1.000 total time=
6.3s
[CV 3/3] END ....max_depth=20, n_estimators=200;; score=1.000 total time=
13.6s
[CV 2/3] END ....max_depth=30, n_estimators=150;; score=1.000 total time=
7.7s
[CV 1/3] END ....max_depth=30, n_estimators=150;; score=1.000 total time=
10.3s
[CV 3/3] END ....max_depth=30, n_estimators=150;; score=1.000 total time=
9.3s
[CV 1/3] END ....max_depth=30, n_estimators=200;; score=1.000 total time=
11.8s
[CV 2/3] END ....max_depth=30, n_estimators=200;; score=1.000 total time=
9.3s
[CV 3/3] END ....max_depth=30, n_estimators=200;; score=1.000 total time=
8.2s

```

```

Out[ ]:
└─ GridSearchCV
  └─ estimator: RandomForestClassifier
    └─ RandomForestClassifier

```

```
In [ ]: grid_rf_clf.best_score_
```

Out[]: 1.0

In []: `grid_rf_clf.best_params_`

Out[]: `{'max_depth': 30, 'n_estimators': 100}`

In []: `optimal_rf_clf = grid_rf_clf.best_estimator_`

In []: `# Get the predicted classes`
`train_class_preds = optimal_rf_clf.predict(X_train)`
`test_class_preds = optimal_rf_clf.predict(X_test)`

In []: `# Get the accuracy scores`
`train_accuracy_rf = accuracy_score(train_class_preds,y_train)`
`test_accuracy_rf = accuracy_score(test_class_preds,y_test)`

`print("The accuracy on train data is ", train_accuracy_rf)`
`print("The accuracy on test data is ", test_accuracy_rf)`

The accuracy on train data is 1.0
 The accuracy on test data is 1.0

In []: `test_accuracy_rf = accuracy_score(test_class_preds,y_test)`
`test_precision_score_rf = precision_score(test_class_preds,y_test)`
`test_recall_score_rf = recall_score(test_class_preds,y_test)`
`test_f1_score_rf = f1_score(test_class_preds,y_test)`
`test_roc_score_rf = roc_auc_score(test_class_preds,y_test)`

`print("The accuracy on test data is ", test_accuracy_rf)`
`print("The precision on test data is ", test_precision_score_rf)`
`print("The recall on test data is ", test_recall_score_rf)`
`print("The f1 on test data is ", test_f1_score_rf)`
`print("The roc_score on test data is ", test_roc_score_rf)`

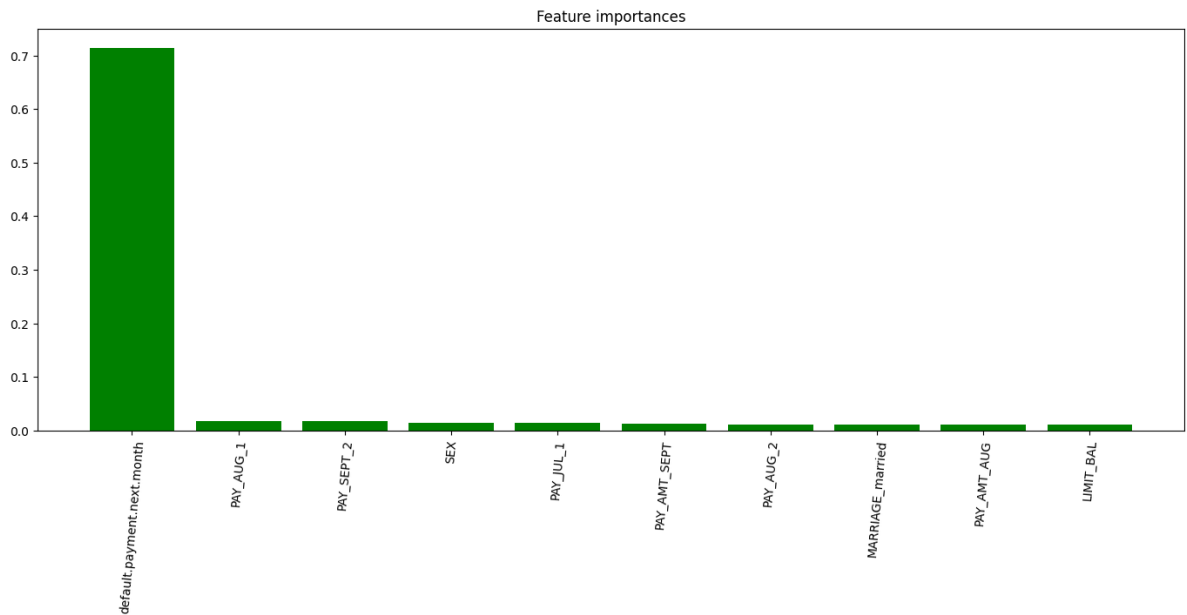
The accuracy on test data is 1.0
 The precision on test data is 1.0
 The recall on test data is 1.0
 The f1 on test data is 1.0
 The roc_score on test data is 1.0

In []: `len(optimal_rf_clf.feature_importances_)`

Out[]: 81

In []: `# Feature Importance`
`feature_importances_rf = pd.DataFrame(optimal_rf_clf.feature_importances_,`
 `index = columns,`
 `columns=['importance_rf']).sort_values(`
 `asc`

`plt.subplots(figsize=(17,6))`
`plt.title("Feature importances")`
`plt.bar(feature_importances_rf.index, feature_importances_rf['importance_rf'`
 `color="g", align="center")`
`plt.xticks(feature_importances_rf.index, rotation = 85)`
`#plt.xlim([-1, X.shape[1]])`
`plt.show()`



```
In [ ]: model_save_name = 'rf_optimized_classifier.pt'
        path = F"./{model_save_name}"
        torch.save(optimal_rf_clf, path)
```


```
In [ ]: model_save_name = 'rf_optimized_classifier.pt'
        path = F"./{model_save_name}"
        optimal_rf_clf = torch.load(path)
```

```
In [ ]: # Get the predicted classes
        train_class_preds = optimal_rf_clf.predict(X_train)
        test_class_preds = optimal_rf_clf.predict(X_test)
```

```
In [ ]: y_preds_proba_rf = optimal_rf_clf.predict_proba(X_test)[::,1]
```

```
In [ ]: import sklearn.metrics as metrics
```

```
In [ ]: y_pred_proba = y_preds_proba_rf
        fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
        auc = metrics.roc_auc_score(y_test, y_pred_proba)
        plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
        plt.legend(loc=4)
        plt.show()
```



Implementing XGBoost

```
In [ ]: #import lightgbm and xgboost
import lightgbm as lgb
import xgboost as xgb
```

Applying XGBoost

```
In [ ]: #The data is stored in a DMatrix object
#label is used to define our outcome variable
dtrain=xgb.DMatrix(X_train,label=y_train)
dtest=xgb.DMatrix(X_test)
```

```
In [ ]: #setting parameters for xgboost
parameters={'max_depth':7, 'eta':1, 'silent':1,'objective':'binary:logistic'}
```

```
In [ ]: #training our model
num_round=50
from datetime import datetime
start = datetime.now()
xg=xgb.train(parameters,dtrain,num_round)
stop = datetime.now()
```

[16:37:36] WARNING: ../src/learner.cc:627:
Parameters: { "silent" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

```
In [ ]: #Execution time of the model
        execution_time_xgb = stop-start
        execution_time_xgb
```

```
Out[ ]: datetime.timedelta(seconds=2, microseconds=326039)
```

```
In [ ]: #now predicting our model on train set
        train_class_preds_probs=xg.predict(dtrain)
        #now predicting our model on test set
        test_class_preds_probs =xg.predict(dtest)
```

```
In [ ]: len(train_class_preds_probs)
```

```
Out[ ]: 31307
```

```
In [ ]: train_class_preds = []
        test_class_preds = []
        for i in range(0,len(train_class_preds_probs)):
            if train_class_preds_probs[i] >= 0.5:
                train_class_preds.append(1)
            else:
                train_class_preds.append(0)

        for i in range(0,len(test_class_preds_probs)):
            if test_class_preds_probs[i] >= 0.5:
                test_class_preds.append(1)
            else:
                test_class_preds.append(0)
```

```
In [ ]: test_class_preds_probs[:20]
```

```
Out[ ]: array([0.04043363, 0.04043363, 0.04043363, 0.04043363, 0.04043363,
               0.04043363, 0.04043363, 0.04043363, 0.95956635, 0.04043363,
               0.95956635, 0.04043363, 0.95956635, 0.04043363, 0.95956635,
               0.04043363, 0.95956635, 0.04043363, 0.95956635, 0.04043363],
              dtype=float32)
```

```
In [ ]: test_class_preds[:20]
```

```
Out[ ]: [0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1]
```

```
In [ ]: len(y_train)
```

```
Out[ ]: 31307
```

```
In [ ]: len(train_class_preds)
```

```
Out[ ]: 31307
```

```
In [ ]: # Get the accuracy scores
train_accuracy_xgb = accuracy_score(train_class_preds,y_train)
test_accuracy_xgb = accuracy_score(test_class_preds,y_test)

print("The accuracy on train data is ", train_accuracy_xgb)
print("The accuracy on test data is ", test_accuracy_xgb)
```

The accuracy on train data is 1.0

The accuracy on test data is 1.0

```
In [ ]: test_accuracy_xgb = accuracy_score(test_class_preds,y_test)
test_precision_xgb = precision_score(test_class_preds,y_test)
test_recall_score_xgb = recall_score(test_class_preds,y_test)
test_f1_score_xgb = f1_score(test_class_preds,y_test)
test_roc_score_xgb = roc_auc_score(test_class_preds,y_test)

print("The accuracy on test data is ", test_accuracy_xgb)
print("The precision on test data is ", test_precision_xgb)
print("The recall on test data is ", test_recall_score_xgb)
print("The f1 on test data is ", test_f1_score_xgb)
print("The roc_score on train data is ", test_roc_score_xgb)
```

The accuracy on test data is 1.0

The precision on test data is 1.0

The recall on test data is 1.0

The f1 on test data is 1.0

The roc_score on train data is 1.0

Hyperparameter Tuning

```
In [ ]: from xgboost import XGBClassifier
```

```
In [ ]: X = df_fr.drop(['IsDefaulter','Payement_Value','Dues'],axis=1)
y = df_fr['IsDefaulter']
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, r
```

```
In [ ]: param_test1 = {
    'max_depth':range(3,10,2),
    'min_child_weight':range(1,6,2)
}
gsearch1 = GridSearchCV(estimator = XGBClassifier( learning_rate =0.1, n_es
min_child_weight=1, gamma=0, subsample=0.8, colsample_bytree=0.8,
objective= 'binary:logistic', nthread=4, scale_pos_weight=1, seed=27),
    param_grid = param_test1, scoring='accuracy',n_jobs=-1, cv=3, verbose = 2)
gsearch1.fit(X_train, y_train)
```

Fitting 3 folds for each of 12 candidates, totalling 36 fits


```

Out[ ]: GridSearchCV(cv=3,
                    estimator=XGBClassifier(base_score=None, booster=None,
                                           callbacks=None, colsample_bylevel=None,
                                           colsample_bynode=None,
                                           colsample_bytree=0.8,
                                           early_stopping_rounds=None,
                                           enable_categorical=False, eval_metric=
None,
                                           gamma=0, gpu_id=None, grow_policy=None,
                                           importance_type=None,
                                           interaction_constraints=None,
                                           learning_rate=0.1, max_bin=None,
                                           max_cat_to_onehot=None,
                                           max_delta_step=None, max_depth=5,
                                           max_leaves=None, min_child_weight=1,
                                           missing=nan, monotone_constraints=None,
                                           n_estimators=140, n_jobs=None, nthread
=4,
                                           num_parallel_tree=None, predictor=None,
                                           random_state=None, reg_alpha=None,
                                           ...),
                    n_jobs=-1,
                    param_grid={'max_depth': range(3, 10, 2),
                                'min_child_weight': range(1, 6, 2)},
                    scoring='accuracy', verbose=2)

```

```
In [ ]: gsearch1.best_score_
```

```
Out[ ]: 1.0
```

```
In [ ]: optimal_xgb = gsearch1.best_estimator_
```

```
In [ ]: # Get the predicted classes
train_class_preds = optimal_xgb.predict(X_train)
test_class_preds = optimal_xgb.predict(X_test)
```

```
In [ ]: # Get the accuracy scores
train_accuracy_xgb_tuned = accuracy_score(train_class_preds,y_train)
test_accuracy_xgb_tuned = accuracy_score(test_class_preds,y_test)

print("The accuracy on train data is ", train_accuracy_xgb_tuned)
print("The accuracy on test data is ", test_accuracy_xgb_tuned)
```

```

The accuracy on train data is  1.0
The accuracy on test data is  1.0

```

```
In [ ]: test_accuracy_xgb_tuned = accuracy_score(test_class_preds,y_test)
test_precision_xgb_tuned = precision_score(test_class_preds,y_test)
test_recall_score_xgb_tuned = recall_score(test_class_preds,y_test)
test_f1_score_xgb_tuned = f1_score(test_class_preds,y_test)
test_roc_score_xgb_tuned = roc_auc_score(test_class_preds,y_test)

print("The accuracy on test data is ", test_accuracy_xgb_tuned)
print("The precision on test data is ", test_precision_xgb_tuned)
print("The recall on test data is ", test_recall_score_xgb_tuned)
print("The f1 on test data is ", test_f1_score_xgb_tuned)
print("The roc_score on train data is ", test_roc_score_xgb_tuned)
```

The accuracy on test data is 1.0
The precision on test data is 1.0
The recall on test data is 1.0
The f1 on test data is 1.0
The roc_score on train data is 1.0


```
In [ ]: pd.DataFrame(optimal_xgb.feature_importances_,
                    index = columns,
                    columns=['importance_xgb']).sort_values
                    asc
```

Out[]:

	importance_xgb
default.payment.next.month	0.774541
PAY_AUG_1	0.061343
PAY_SEPT_2	0.058241
PAY_JUL_1	0.036967
SEX	0.016271
PAY_AUG_2	0.014785
MARRIAGE_married	0.008888
PAY_JUN_1	0.006657
PAY_JUL_-1	0.005823
PAY_SEPT_1	0.005516


```
In [ ]: # Feature Importance
feature_importances_xgb = pd.DataFrame(optimal_xgb.feature_importances_,
                                       index = columns,
                                       columns=['importance_xgb']).sort_values
                                       asc

plt.subplots(figsize=(17,6))
plt.title("Feature importances")
plt.bar(feature_importances_xgb.index, feature_importances_xgb['importance_
        color="b", align="center")
plt.xticks(feature_importances_rf.index, rotation = 85)
#plt.xlim([-1, X.shape[1]])
plt.show()
```



```
In [ ]: y_preds_proba_xgb = optimal_xgb.predict_proba(X_test)[::,1]
```

```
In [ ]: y_pred_proba = y_preds_proba_xgb
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr, tpr, label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



```
In [ ]: model_save_name = 'xgb_optimized_classifier.pt'
path = F"./{model_save_name}"
torch.save(optimal_xgb, path)
```

```
In [ ]: model_save_name = 'xgb_optimized_classifier.pt'
path = F"./{model_save_name}"
optimal_xgb = torch.load(path)
```

Evaluating the models

```
In [ ]: recall_score
```

```
Out[ ]: <function sklearn.metrics._classification.recall_score(y_true, y_pred, *, l
labels=None, pos_label=1, average='binary', sample_weight=None, zero_divisio
n='warn')>
```

```
In [ ]: classifiers = ['Logistic Regression', 'SVC', 'Random Forest CLf', 'Xgboost
train_accuracy = [train_accuracy_lr, train_accuracy_SVC, train_accuracy_rf,
test_accuracy = [test_accuracy_lr, test_accuracy_SVC, test_accuracy_rf, tes
precision_score = [test_precision_score_lr, test_precision_score_SVC, test_
recall_score = [test_recall_score_lr, test_recall_score_SVC, test_recall_sc
f1_score = [test_f1_score_lr, test_f1_score_SVC, test_f1_score_rf, test_f1_
```

```
In [ ]: pd.DataFrame({'Classifier':classifiers, 'Train Accuracy': train_accuracy, '

```

Out[]:

	Classifier	Train Accuracy	Test Accuracy	Precision Score	Recall Score	F1 Score
0	Logistic Regression	1.0	1.000000	1.000000	1.000000	1.0000
1	SVC	1.0	0.996693	0.998962	0.994448	0.9967
2	Random Forest Clf	1.0	1.000000	1.000000	1.000000	1.0000
3	Xgboost Clf	1.0	1.000000	1.000000	1.000000	1.0000

Plotting ROC AUC for all the models

```
In [ ]: classifiers_proba = [(optimized_clf, y_preds_proba_lr),
                             (optimal_rf_clf, y_preds_proba_rf),
                             (optimal_xgb, y_preds_proba_xgb),
                             (optimal_SVC_clf,y_pred_proba_SVC)]

# Define a result table as a DataFrame
result_table = pd.DataFrame(columns=['classifiers', 'fpr','tpr','auc'])

# Train the models and record the results
for pair in classifiers_proba:

    fpr, tpr, _ = roc_curve(y_test, pair[1])
    auc = roc_auc_score(y_test, pair[1])

    result_table = result_table.append({'classifiers':pair[0].__class__.__name__,
                                       'fpr':fpr,
                                       'tpr':tpr,
                                       'auc':auc}, ignore_index=True)

# Set name of the classifiers as index labels
result_table.set_index('classifiers', inplace=True)
```

In []: result_table

Out[]:

	fpr	tpr	auc
classifiers			
LogisticRegression	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	[0.0, 0.00012970168612191958, 0.15642023346303...	1.000000
RandomForestClassifier	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	[0.0, 0.0009079118028534371, 0.001297016861219...	1.000000
XGBClassifier	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	[0.0, 0.0031128404669260703, 0.004150453955901...	1.000000
SVC	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	[0.0, 0.00012970168612191958, 0.21089494163424...	0.999977

```
In [ ]: fig = plt.figure(figsize=(8,6))

for i in result_table.index:
    plt.plot(result_table.loc[i]['fpr'],
             result_table.loc[i]['tpr'],
             label="{}, AUC={:.3f}".format(i, result_table.loc[i]['auc']))

plt.plot([0,1], [0,1], color='orange', linestyle='--')
```

```
plt.xticks(np.arange(0.0, 1.1, step=0.1))
plt.xlabel("Flase Positive Rate", fontsize=15)

plt.yticks(np.arange(0.0, 1.1, step=0.1))
plt.ylabel("True Positive Rate", fontsize=15)

plt.title('ROC Curve Analysis', fontweight='bold', fontsize=15)
plt.legend(prop={'size':13}, loc='lower right')

plt.show()
```

